

1 Linked-List

The original algorithm by Harris is presented in Figure 1. Harris approach uses an Atomic-Markable-Reference object, in which the next field of a Node, in addition to a reference to the next node in the list, is also marked or unmarked. The two fields can be update atomically, either together or individually. This can be done by using the most-significant-bit of next for the marking. For simplicity, we assume `node.next` returns the reference, while a query can be used to identify if it is marked. Therefore, whenever writing to `node.next`, or performing CAS, both the reference and marking state should be mention.

For ease of presentation, we assume a List is initialised with head and tail, containing keys $\infty, -\infty$ respectively, while we do not allow insert or delete of these keys.

The Lookup procedure is used by Insert and Delete, in order to find the node with the lowest key greater or equal to the input key, and its predecessor on the list, while physically removing any marked node on its way. To insert a key α , a process first finds the right location for α using the Lookup procedure, and then tries to set `pred.next` to point to a new node containing α by performing CAS. To delete a key α , a process looks for it using the Lookup procedure, and then tries to logically remove it by marking `curr.next` using CAS. In case the marking was successful, the process also tries to physically remove the node. To find a key α , a process simply looks to see if there is a node in the list with key α which is unmarked.

Procedure Lookup(int key)

Data: Node* pred, curr, succ

```
1 retry: while true do
2   | pred = head
3   | curr = head.next
4   | while true do
5   |   | succ = curr.next
6   |   | if curr.next is marked then
7   |   |   | if pred.next.CAS (unmarked curr, unmarked succ) == false then
8   |   |   |   | go to retry
9   |   |   | end
10  |   |   | curr = succ
11  |   | else
12  |   |   | if curr.key ≥ key then
13  |   |   |   | return <pred, curr>
14  |   |   | end
15  |   |   | pred = curr
16  |   |   | curr = succ
17  |   | end
18  | end
19 end
```

Shared variables: Node* head

Procedure Insert(int key)	
<hr/>	
Data: Node* pred, curr	
Node node = new Node (key)	
20	while true do
21	<pred, curr> = lookup(key)
22	if curr.key == key then
23	return false
24	else
25	node.next = unmarked curr
26	if pred.next.CAS (unmarked curr, unmarked node) then
27	return true
28	end
29	end
30	end
<hr/>	
Procedure Delete(int key)	
<hr/>	
Data: Node* pred, curr, succ	
31	while true do
32	<pred, curr> = lookup(key)
33	if curr.key != key then
34	return false
35	else
36	succ = curr.next
37	if curr.next.CAS (unmarked succ, marked succ) then
38	pred.next.CAS (unmarked curr, unmarked succ)
39	return true
40	end
41	end
42	end
<hr/>	
Procedure Find(int key)	
<hr/>	
Data: Node* curr = head	
43	while curr.key < key do
44	curr = curr.next
45	end
46	return (curr.key == key && curr.next is unmarked)

Figure 1: Harris Non-Blocking Algorithm

1.0.1 Crash-Recovery

The linearization point are as follows:

Insert: At the point of a successful CAS

Delete: At the point of a successful CAS for marking the node (logical delete)

Find: At the point of the procedure return, that is, either when `curr.key != key`, or at the second condition test.

Following these linearization points (committing proof...), insert and delete operation are linearized at the point where they affect the system. That is, if there is a linearization point to insert operation, then all process will see the new node starting from this point, and if a node was logically removed, then all processes treat it as a removed node. Therefore, once a process recovers following a crash, the List data structure is consistent - if it has a pending operation, then either this operation already had a linearization point and affect all other processes, or that it did not affect the data structure at all.