

# 1 Robust BST

The original BST algorithm does not support the crash-recovery model. It is clear from the code a process does not persist the operation's response in the non-volatile memory, and thus, once a process crash the response is lost. For example, assume a process  $q$  apply  $\text{INSERT}(k)$ , performs a successful CAS in line 101 and fails after completing the  $\text{HELPINSERT}$  routine. In this case, the  $\text{INSERT}$  operation took effect, that is, the new key appears as a leaf in the tree, and any  $\text{FIND}(k)$  operation will return it. However, even though the operation must be linearized before the crash, upon recovery process  $q$  is unaware of it. Moreover, looking for the new leaf in the tree may be futile, as it might be  $k$  has been removed from the tree after the crash.

Furthermore, if no recover routine is supplied, it may result an execution which is not well-formed. Consider for example the following scenario. A process  $q$  invoke an  $Op_1 = \text{INSERT}(k_1)$  operation.  $q$  performs a successful CAS in line 101 followed by a crash. After recovering,  $q$  invoke an  $Op_2 = \text{INSERT}(k_2)$  operation. Assume  $k_1$  and  $k_2$  belongs to a different parts of the tree (do not share parent or grandparent). Then,  $q$  can complete the insertion of  $k_2$  without having any affect on  $k_1$ . Now, a process  $q'$  performs  $\text{FIND}(k_1)$  which returns  $\perp$ , as the insertion of  $k_1$  is not completed, followed by  $\text{FIND}(k_2)$ , which returns the leaf of  $k_2$ . The  $\text{INSERT}(k_1)$  operation will be completed later by any  $\text{INSERT}$  or  $\text{DELETE}$  operation which needs to make changes to the flagged node. We get that  $Op_2$  must be before  $Op_1$  in the linearization, although  $Op_1$  invoked first.

The kind of anomaly described above can be addressed by having the first CAS of a successful attempt for  $\text{INSERT}$  or  $\text{DELETE}$  as the linearization point, as in the Linked-List. For that, the  $\text{FIND}$  routine should take into consideration future unavoidable changes, for example, a node flagged with  $\text{IFlag}$  ensures an insertion of some key. A simple solution is to change the  $\text{FIND}$  routine such that it also helps other operations, as described in figure 1. The  $\text{FIND}$  routine will search for key  $k$  in the tree. If the  $\text{SEARCH}$  routine returns a grandparent or a parent that is flagged, then it might be that an insert or delete of  $k$  is currently in progress, thus we first help the operation to complete, and then search for  $k$  again. Otherwise, if  $gpupdate$  or  $pupdate$  has been changed since the last read, it means some change already took affect, and there is a need to search for  $k$  again. If none of the above holds, there is a point in time where  $gp$  points to  $p$  which points to  $l$ , and there is no attempt to change this part of the tree. As a result, if  $k$  is in the tree at this point, it must be in  $l$ , and the find can return safely.

The approach described above is not efficient in terms of time. We would like a solution which maintain the desirable behaviour of the original  $\text{FIND}$  routine, where a single  $\text{SEARCH}$  is needed. A more refined solution is given in figure 2. The intuition for it is drawn from the Linked-List algorithm. In the Linked-List algorithm it was enough to consider a marked node as if it has been deleted, without the need to complete the deletion. Nonetheless, the complex BST implementation is more challenging, as the  $\text{DELETE}$  routine needs to successfully capture two nodes using CAS in order to complete the deletion. Therefore, if a process  $p$  executes  $\text{FIND}(k)$  procedure, and observes a node flagged with  $\text{DFlag}$  attempting to delete the key  $k$ , it can not know whether in the future this delete attempt will succeed or fail, and thus does not know whether to consider the key  $k$  as part of the tree or not. To overcome this problem, in such case the process will first try and validate the delete operation by marking the relevant node. According to whether the marking attempt was successful, the process can conclude if the delete operation is successful or not. In order to easily implement the modified  $\text{FIND}$  routine there is a need to conclude from  $\text{IInfo}$  what is the new leaf (leaf *new* in the  $\text{INSERT}$  routine). For simplicity of presentation, we do not add this field, and abstractly refer to it in the code.

The correctness of the two suggested solutions relies on the following argument. Once a process flags a node during operation  $Op$  with input key  $k$  (either INSERT or DELETE), then if this attempt to complete the operation eventually succeed (i.e., the marking is also successful in the case of DELETE), then any FIND( $k$ ) operation invoked from this point consider  $Op$  as if it is completed.

The suggested modification, although being simple and local, only guarantee the implementation satisfy R-linearizability. However, the problem of response being lost in case of a crash is not addressed. Roughly speaking, the critical points in the code for recovery are the CAS primitives, as a crash right after applying CAS operation results the lost of the response, and in order to complete the operation the process needs to know the result of the CAS. In addition, because of the helping mechanism, a suspended DELETE operation which flagged a node and yet to mark one, may be completed by other process in the future, and may not. Upon recovery, the process needs to distinguish between the two cases, in order to obtain the right response.

To address this issue, we expend the helping mechanism so that it also update the info structure in case of a success. This is done by adding a boolean field, *done*, to the Flag structure. This way, if a process crash along an operation  $Op$ , upon recovery it can check to see if the operation was already completed. A crucial point is to update the *done* bit before performing the unflagging. Therefore, if a node is no longer flagged we can be sure *done* was already updated. If we switch the order, then it might be an operation and unflagging were completed, but the *done* bit is yet to be updated. Therefore, other processes can change the BST structure. However, if the process crash and recover at this point, the *done* bit is off, and the BST structure has been changed, so it will be harder for the process to conclude whether the operation took affect.

Before a process  $q$  attempt to perform an operation, as it creates the Flag structure  $op$  describing the operation and its affect on the data structure, the process stores  $op$  in a designated location in the shared memory (for simplicity, we use an array). As a result, upon recovery  $q$  has an access to this information. Now,  $q$  can check to see if the operation is still in progress, i.e., if the relevant node (parent or grandparent) is still flagged. If so, it first tries to complete the operation. Otherwise, it implies either the operation was completed, and therefore *done* bit is updated, or that the attempt was unsuccessful and there was no write to the *done* bit. Hence, the *done* bit can distinguish between the two scenarios. Notice that there is a scenario in which process  $q$  recovers and observes an operation  $Op$  as it in a progress, but just before it retries it, some other process complete the operation. We need to prove that even in such case, the operation will affect the data structure exactly once, and the right response is returned.

The given implementation does not recover the FIND routine, since this routine does not make any changes to the BST, hence it is always safe to consider it as having no linearization point and reissue it. Also, for ease of presentation, we only write to  $Announce[id]$  once we are about to capture a node using a CAS. However, writing to  $Announce[id]$  at the beginning of the routine may be helpful in case of a crash early in the routine, so that the process will be able to use the data stored in  $Announce[id]$  in such case also. The same is true with response value,  $Announce[id] \rightarrow done$  is updated only if the routine made changes to the BST.

### 1.0.1 Correctness

In the following section we give a proof sketch for the algorithm correctness. We assume for simplicity nodes and Flag records are always allocated new memory locations, although it is enough to require no location is reallocated as long as there is a chain of pointers leading to it. The proof relies on the correctness of the original algorithm, which can be found on [...].

The proof relies on several key arguments given below.

[Arg1] The original algorithm is anonymous and uniform, i.e., any number of processes can use the BST, and there is no need to know the number of processes in the system in order to use the BST. Notice that all helping routines in the given implementation are completely anonymous, and an execution of such a routine by either the process which invoked  $op$  or any other helping process executes the exact same code. This observation allows the use of the following argument. If a process crash while executing some helping routine, we can consider it as an helping process which stop taking steps (more formally, there is an equivalent execution in which there is such a process, and it is indistinguishable to all process in the system). Since such process can not cause a wrong behaviour of the algorithm, so does the crash. A corollary of this argument is that repeating an helping routine multiple times by the same process can not violate the BST specification, as there is an equivalent executions in which multiple processes executes the different helping routines.

[Arg2] It is easy to verify the post-conditions of the SEARCH routine still holds, as they follow directly from the routine's code, and does not rely on the structure or correctness of the BST. Also, the SEARCH routine does not make any changes to the BST, but rather simply traverse it. Therefore FIND routine, which only uses SEARCH, does not affect any process, and in case of a failure along FIND execution, reissuing it satisfies NRL.

[Arg3] If an internal node  $nd_1$  stops pointing to a node  $nd_2$  at some point of the execution, it can not point to  $nd_2$  again. This attributes to the fact an INSERT presents a node with two new children. Therefore, if  $nd_2$  is a leaf, it can either be delete, or replaced by a new copy of an INSERT operation. Otherwise,  $nd_2$  is an internal node, and as such, the pointer to it by  $nd_1$  can not be replaced by an INSERT operation (which only allows to replacement a leaf), and therefore it can only be removed from the tree.

[Arg4] The field update of a node  $nd$  can have any value only once along an execution. Any attempt to perform an operation creates a new record in the memory. If  $nd \rightarrow update$  is marked, it can not be unmarked or changed. Otherwise, any attempt to flag it uses a new created record  $op$ . If the attempt succeed, then eventually it will be unflagged while still referring to  $op$ . In order to replace the value again, there must be an operation reading  $nd \rightarrow update$  after it was unflagged (as any operation first help a flagged node). This operation must create a new record, and thus we can use the same argument again. As a corollary, if a process successfully flag or mark a node, there was no change to the node since the last time it read the update field of the node.

**Proof Sketch** Assume a process  $q$  performs an operation  $Op$  (either INSERT or DELETE). If  $q$  does not crash, the algorithm is identical to the original algorithm, except for the additional write to  $Announce[q]$  and  $op \rightarrow done$ , and thus the correctness of the original algorithm can be applied. Otherwise,  $q$  crash at some point, and upon recovery it reads  $op$  from  $Announce[q]$ . This record represent the last attempt of  $q$  to complete  $Op$ . We split the proof based on the type of operation.

$Op = \text{INSERT}$ . Consider the read of  $op \rightarrow p \rightarrow update$  upon recovery, and denote this value by  $pupdate$ . If  $pupdate = \langle IFlag, op \rangle$ , this implies the iflag CAS in line 101 was successful and the operation is yet to complete. It might be that INSERT already took affect, that is, the new key is part of the tree, but the unflagging is yet to happen. In such case,  $q$  calls  $\text{HELPINSERT}(op)$  in order to try and complete the operation. Considering arg1, this call can not violate the BST correctness, even if it not the first time  $q$  executes it. Moreover, during  $\text{HELPINSERT}$  there is a write to  $op \rightarrow done$ , and thus after completing the routine  $q$  returns TRUE, as required.

Else  $pupdate \neq \langle IFlag, op \rangle$ . There are two scenarios to consider. Either the iflag CAS of  $q$  in

line 101 was successful or not. If it was successful, then  $p \rightarrow update = \langle IFlag, op \rangle$  at this point. The only way to change it is to first unflag  $p$ . To do so, a process needs to complete an  $HELPINSERT(op)$  routine, and in particular must write to  $op \rightarrow done$ . In such case, the  $INSERT$  operation was completed, and  $q$  returns  $TRUE$ . Otherwise, the CAS was not successful, either because it failed, or the crash was before the CAS. In both cases, the  $INSERT$  operation will not be completed, as  $op$  is not stored in  $p \rightarrow update$ , and thus no process has an access to it. Consequently, no process can update  $op \rightarrow done$ , and  $q$  returns  $FAIL$ .

$Op = DELETE$ . Consider the read of  $op \rightarrow gp \rightarrow update$  upon recovery, and denote this value by  $gpupdate$ . If  $gpupdate = \langle DFlag, op \rangle$ , this implies the  $dflag$  CAS of  $q$  in line 128 was successful, and the operation is yet to complete. As in the  $INSERT$ , it might be the operation already changed the tree. After reading  $gpupdate$   $q$  invokes  $HELPDELETE(op)$  routine. Again, following  $arg1$ , executing this multiple times by  $q$  can not violate the BST correctness. The first process to try and mark  $op \rightarrow p \rightarrow update$  during an  $HELPDELETE(op)$  routine is the one to determine the outcome of it. If it is successful, then  $p$  is marked, and the  $update$  field can not be changed. That is, any  $HELPDELETE(op)$  execution will obtain true in line 139, and will call  $HELPMARKED(op)$  routine. Otherwise, the CAS fails, and so  $p \rightarrow update$  is no longer equal to  $op \rightarrow pupdate$ . By  $arg4$  it will never get this value again, and thus any marking CAS during a  $HELPDELETE(op)$  execution will fail, and there is no call to  $HELPMARKED(op)$ . In the first case, any  $HELPDELETE(op)$  routine must first complete a  $HELPMARKED(op)$ , and thus must write to  $op \rightarrow done$ , while in the later case, there is no write to  $op \rightarrow done$ , as no  $HELPMARKED(op)$  is ever invoked. Therefore, in both cases, when  $q$  completes  $HELPMARKED(op)$  it reads  $op \rightarrow done$  and returns the right response.

Otherwise  $gpupdate \neq \langle DFlag, op \rangle$ , and there are two scenarios to consider. If the  $dflag$  CAS of  $q$  in line 128 never took affect, because it either failed, or the crash preceded it, then  $op$  is never written to  $gp \rightarrow update$ , or to any update field. Thus, no process is aware of it, and  $op \rightarrow done$  remains  $FALSE$ , resulting  $q$  returning  $FAIL$  as required. Else, the CAS was successful, and  $gp \rightarrow update$  was flagged. The only way to change it is to first unflag it, and this in turn can be done only during an  $HELPDELETE(op)$  routine. In this case, it can be unflagged in either the  $HELPMARKED$  routine in line 154, or in line 145 of the  $HELPDELETE$  routine. As mention before, the first CAS in line 138 of an  $HELPDELETE(op)$  execution determines the outcome for all  $HELPDELETE(op)$ . If it is successful,  $p \rightarrow update$  is forever marked, and all  $HELPDELETE(op)$  must invoke  $HELPMARKED(op)$ . Therefore, the only option to unflag  $gp \rightarrow update$  is at the end of  $HELPMARKED(op)$  routine, and this done only after setting  $op \rightarrow done$ . In such case, the  $DELETE$  operation took affect, and  $q$  will return  $TRUE$ . On the other hand, if the CAS was not successful, then any  $HELPDELETE(op)$  will fail to mark  $p \rightarrow update$ , and hence no  $HELPMARKED(op)$  is ever invoked. As a result, there is no write to  $op \rightarrow done$ . In such case, the  $DELETE$  operation did not took affect, nor will be, and indeed  $q$  will return  $FAIL$ .

```

FIND(Key k) : Leaf* {
1      Internal *gp, *p
2      Leaf *l
3      Update pupdate, gpupdate

4      while (TRUE) {
5           $\langle gp, p, l, pupdate, gpupdate \rangle := \text{SEARCH}(k)$ 
6          if gpupdate.state  $\neq$  CLEAN then HELP(gpupdate)
7          else if pupdate.state  $\neq$  CLEAN then HELP(pupdate)
8          else if gp  $\rightarrow$  update = gpupdate and p  $\rightarrow$  update = pupdate then {
9              if l  $\rightarrow$  key = k then return l
10             else return  $\perp$ 
11         }
12     }
13 }

```

Figure 1: Solution 1: R-linearizable FIND routine

```

FIND(Key k) : Leaf* {
14     Internal *gp, *p
15     Leaf *l
16     Update pupdate, gpupdate

17      $\langle gp, p, l, pupdate, gpupdate \rangle := \text{SEARCH}(k)$ 
18     if l  $\rightarrow$  key  $\neq$  k then {
19         if (pupdate.state = IFlag and pupdate.info attempt to add key k) then
20             return leaf with key k from pupdate.info
21         else return  $\perp$ 
22     }
23     if (pupdate.state = MARK and pupdate.info  $\rightarrow$  l  $\rightarrow$  key = k) then return  $\perp$ 
24     if (gpupdate.state = DFlag and gpupdate.info  $\rightarrow$  l  $\rightarrow$  key = k) then {
25         op := gpupdate.info
26         result := CAS(op  $\rightarrow$  p  $\rightarrow$  update, op  $\rightarrow$  pupdate,  $\langle \text{MARK}, op \rangle$ )       $\triangleright$  mark CAS
27         if (result = op  $\rightarrow$  pupdate or result =  $\langle \text{MARK}, op \rangle$ ) then return  $\perp$        $\triangleright$  op  $\rightarrow$  p is successfully marked
28     }
29     return l
30 }

```

Figure 2: Solution 2: R-linearizable FIND routine

```

31 type Update {           ▷ stored in one CAS word
32     {CLEAN, DFlag, IFlag, MARK} state
33     Flag *info
34 }
35 type Internal {         ▷ subtype of Node
36      $Key \cup \{\infty_1, \infty_2\}$  key
37     Update update
38     Node *left, *right
39 }
40 type Leaf {             ▷ subtype of Node
41      $Key \cup \{\infty_1, \infty_2\}$  key
42 }
43 type IInfo {            ▷ subtype of Flag
44     Internal *p, *newInternal
45     Leaf *l
46     boolean done
47 }
48 type DInfo {            ▷ subtype of Flag
49     Internal *gp, *p
50     Leaf *l
51     Update pupdate
52     boolean done
53 }
    ▷ Initialization:
54 shared Internal *Root := pointer to new Internal node
    with key field  $\infty_2$ , update field  $\langle \text{CLEAN}, \perp \rangle$ , and
    pointers to new Leaf nodes with keys  $\infty_1$  and
     $\infty_2$ , respectively, as left and right fields.

```

Figure 3: Type definitions and initialization.

```

RECOVER() {
55     Flag *op = Announce[id]
56     if op of type IInfo then {
57         result := op → p → update           ▷ Check flag
58         if result =  $\langle \text{IFlag}, op \rangle$  then HELPINSERT(op)   ▷ Finish the insertion
59     }
60     if op of type DInfo then {
61         result := op → gp → update           ▷ Check flag
62         if result =  $\langle \text{DFlag}, op \rangle$  then HELPDELETE(op)   ▷ Either finish deletion or unflag
63     }
64     if op → done = TRUE then return TRUE
65     else return FAIL
66 }

```

Figure 4: RECOVER routine

```

67 SEARCH(Key k) : (Internal*, Internal*, Leaf*, Update, Update) {
    ▷ Used by INSERT, DELETE and FIND to traverse a branch of the BST; satisfies following postconditions:
    ▷ (1)  $l$  points to a Leaf node and  $p$  points to an Internal node
    ▷ (2) Either  $p \rightarrow left$  has contained  $l$  (if  $k < p \rightarrow key$ ) or  $p \rightarrow right$  has contained  $l$  (if  $k \geq p \rightarrow key$ )
    ▷ (3)  $p \rightarrow update$  has contained  $pupdate$ 
    ▷ (4) if  $l \rightarrow key \neq \infty_1$ , then the following three statements hold:
    ▷ (4a)  $gp$  points to an Internal node
    ▷ (4b) either  $gp \rightarrow left$  has contained  $p$  (if  $k < gp \rightarrow key$ ) or  $gp \rightarrow right$  has contained  $p$  (if  $k \geq gp \rightarrow key$ )
    ▷ (4c)  $gp \rightarrow update$  has contained  $gpupdate$ 
68   Internal *gp, *p
69   Node *l := Root
70   Update gpupdate, pupdate
    ▷ Each stores a copy of an update field
71   while  $l$  points to an internal node {
72       gp := p
73       p := l
74       gpupdate := pupdate
75       pupdate := p → update
76       if  $k < l \rightarrow key$  then  $l := p \rightarrow left$  else  $l := p \rightarrow right$ 
77   }
78   return (gp, p, l, pupdate, gpupdate)
79 }

80 FIND(Key k) : Leaf* {
81   Leaf *l
82   (−, −, l, −, −) := SEARCH(k)
83   if  $l \rightarrow key = k$  then return  $l$ 
84   else return  $\perp$ 
85 }

86 INSERT(Key k) : boolean {
87   Internal *p, *newInternal
88   Leaf *l, *newSibling
89   Leaf *new := pointer to a new Leaf node whose key field is  $k$ 
90   Update pupdate, result
91   IInfo *op
92   while TRUE {
93       (−, p, l, pupdate, −) := SEARCH(k)
94       if  $l \rightarrow key = k$  then return FALSE
95       if  $pupdate.state \neq CLEAN$  then HELP(pupdate)
96       else {
97         newSibling := pointer to a new Leaf whose key is  $l \rightarrow key$ 
98         newInternal := pointer to a new Internal node with key field  $\max(k, l \rightarrow key)$ ,
          update field (CLEAN,  $\perp$ ), and with two child fields equal to new and newSibling
          (the one with the smaller key is the left child)
99         op := pointer to a new IInfo record containing (p, l, newInternal, FALSE)
100        Announce[id] := op
101        result := CAS(p → update, pupdate, (IFlag, op))
102        if result = pupdate then {
103            HELPINSERT(op)
104            return TRUE
105        }
106        else HELP(result)
107      }
108   }
109 }

110 HELPINSERT(IInfo *op) {
    ▷ Precondition:  $op$  points to an IInfo record (i.e., it is not  $\perp$ )
111   CAS-CHILD(op → p, op → l, op → newInternal)
112   op → done := TRUE
113   CAS(op → p → update, (IFlag, op), (CLEAN, op))
114 }
    ▷ ichild CAS
    ▷ announce the operation completed
    ▷ iunflag CAS

```

Figure 5: Pseudocode for SEARCH, FIND and INSERT.

```

115 DELETE(Key  $k$ ) : boolean {
116   Internal * $gp$ , * $p$ 
117   Leaf * $l$ 
118   Update  $pupdate$ ,  $gpupdate$ ,  $result$ 
119   DInfo * $op$ 

120   while TRUE {
121      $\langle gp, p, l, pupdate, gpupdate \rangle := \text{SEARCH}(k)$ 
122     if  $l \rightarrow key \neq k$  then return FALSE ▷ Key  $k$  is not in the tree
123     if  $gpupdate.state \neq \text{CLEAN}$  then HELP( $gpupdate$ )
124     else if  $pupdate.state \neq \text{CLEAN}$  then HELP( $pupdate$ )
125     else { ▷ Try to flag  $gp$ 
126        $op :=$  pointer to a new DInfo record containing  $\langle gp, p, l, pupdate, \text{FALSE} \rangle$ 
127        $\text{Announce}[id] := op$ 
128        $result := \text{CAS}(gp \rightarrow update, gpupdate, \langle \text{DFlag}, op \rangle)$  ▷ dflag CAS
129       if  $result = gpupdate$  then { ▷ CAS successful
130         if HELPDELETE( $op$ ) then return TRUE ▷ Either finish deletion or unflag
131       }
132       else HELP( $result$ ) ▷ The dflag CAS failed; help the operation that caused the failure
133     }
134   }
135 }

136 HELPDELETE(DInfo * $op$ ) : boolean {
137   ▷ Precondition:  $op$  points to a DInfo record (i.e., it is not  $\perp$ )
138   Update  $result$  ▷ Stores result of mark CAS
139    $result := \text{CAS}(op \rightarrow p \rightarrow update, op \rightarrow pupdate, \langle \text{MARK}, op \rangle)$  ▷ mark CAS
140   if  $result = op \rightarrow pupdate$  or  $result = \langle \text{MARK}, op \rangle$  then { ▷  $op \rightarrow p$  is successfully marked
141     HELPMARKED( $op$ ) ▷ Complete the deletion
142     return TRUE ▷ Tell DELETE routine it is done
143   }
144   else { ▷ The mark CAS failed
145     HELP( $result$ ) ▷ Help operation that caused failure
146      $\text{CAS}(op \rightarrow gp \rightarrow update, \langle \text{DFlag}, op \rangle, \langle \text{CLEAN}, op \rangle)$  ▷ backtrack CAS
147     return FALSE ▷ Tell DELETE routine to try again
148   }

149 HELPMARKED(DInfo * $op$ ) {
150   ▷ Precondition:  $op$  points to a DInfo record (i.e., it is not  $\perp$ )
151   Node * $other$ 
152   ▷ Set  $other$  to point to the sibling of the node to which  $op \rightarrow l$  points
153   if  $op \rightarrow p \rightarrow right = op \rightarrow l$  then  $other := op \rightarrow p \rightarrow left$  else  $other := op \rightarrow p \rightarrow right$ 
154   ▷ Splice the node to which  $op \rightarrow p$  points out of the tree, replacing it by  $other$ 
155   CAS-CHILD( $op \rightarrow gp, op \rightarrow p, other$ ) ▷ dchild CAS
156    $op \rightarrow done := \text{TRUE}$  ▷ announce the operation completed
157    $\text{CAS}(op \rightarrow gp \rightarrow update, \langle \text{DFlag}, op \rangle, \langle \text{CLEAN}, op \rangle)$  ▷ dunflag CAS
158 }

159 HELP(Update  $u$ ) { ▷ General-purpose helping routine
160   ▷ Precondition:  $u$  has been stored in the  $update$  field of some internal node
161   if  $u.state = \text{IFlag}$  then HELPINSERT( $u.info$ )
162   else if  $u.state = \text{MARK}$  then HELPMARKED( $u.info$ )
163   else if  $u.state = \text{DFlag}$  then HELPDELETE( $u.info$ )
164 }

165 CAS-CHILD(Internal * $parent$ , Node * $old$ , Node * $new$ ) {
166   ▷ Precondition:  $parent$  points to an Internal node and  $new$  points to a Node (i.e., neither is  $\perp$ )
167   ▷ This routine tries to change one of the child fields of the node that  $parent$  points to from  $old$  to  $new$ .
168   if  $new \rightarrow key < parent \rightarrow key$  then
169      $\text{CAS}(parent \rightarrow left, old, new)$ 
170   else
171      $\text{CAS}(parent \rightarrow right, old, new)$ 
172 }

```

Figure 6: Pseudocode for DELETE and some auxiliary routines.