

1 Robust BST

The original BST algorithm does not support the crash-recovery model. It is clear from the code a process does not persist the operation's response in the non-volatile memory, and thus, once a process crash the response is lost. For example, assume a process q apply $\text{INSERT}(k)$ and assume q performs a successful CAS in line 1 and fails after completing the HELPINSERT routine. In this case, the INSERT operation took effect, that is, the new key appears as a leaf of the tree, and any $\text{FIND}(k)$ operation will return it. However, even though the operation was already linearized at the time of the crash, upon recovery process q is unaware of it. Moreover, looking for the new leaf in the tree is not healpful, since it might be k has been removed from the tree after the crash.

Moreover, if no recover routine is supplied, it may result an execution which is not well-formed. Consider for example the following scenario. A process q invoke an $Op_1 = \text{INSERT}(k_1)$ operation. After a successful CAS in line 1 the process crash. After recovering, q invoke an $Op_2 = \text{INSERT}(k_2)$ operation. Assume k_1 and k_2 belongs to a different parts of the tree (do not share parent or grandparent). Then, q can complete inserting k_2 without having any affect on k_1 . Now, a process q' performs $\text{FIND}(k_1)$ which returns \perp , as the insertion of k_1 is not completed, follows by an $\text{FIND}(k_2)$, which returns the leaf of k_2 . The $\text{INSERT}(k_1)$ operation will be completed later by any INSERT or DELETE which needs to make changes to the flagged node. We get that Op_2 must be before Op_1 in the linearization, although Op_1 invoked first.

The kind of anomaly described above can be addressed by having the first CAS of a successful attempt for INSERT or DELETE as the linearization point, as in the Linked-List. For that, the FIND routine should take into consideration future unavoidable changes, for example, a node flagged with IFlag ensures an insertion of some key. A simple solution is to change the FIND routine, such that it also helps other operations, as described in figure 1. The FIND routine will search for key k in the tree. If the SEARCH routine returns grandparent or parent that flagged, then it might be that an insert or delete of k is currently in a progress, thus we first help these operation to complete, and then search for k again. Otherwise, if $gpupdate$ or $pupdate$ has been changed since the last read, it means some change already took affect, and there is a need to search for k again. If none of the above holds, there is a point in time where gp points to p which points to l , and there is no attempt to change this part of the tree. As a result, if k is in the tree at this point, it must be in l , and the find can return safely.

The approach described above is not efficient in terms of time. We would like a solution which maintain the desirable behaviour of the original FIND routine, where a single SEARCH is needed. A more refined solution is given in figure 2. The intuition for it is drawn from the Linked-List algorithm. In the Linked-List algorithm it was enough to consider a marked node as if it has been already deleted, without the need to complete the deletion. Nonetheless, the complex BST implementation is more challenging, as the DELETE routine needs to successfully capture two nodes using CAS in order to complete the deletion. Therefore, if a process p executes $\text{FIND}(k)$ procedure, and observes a node flagged with DFlag attempting to delete the key k , it can not know whether in the future this delete attempt will succeed or fail, and thus does not know whether to consider the key k as part of the tree or not. To overcome this problem, in such case the process will first try and validate the delete operation by marking the relevant node. According to whether the marking attempt was successful, the process can conclude if the delete operation is successful or not. In order to easily implement the modified FIND routine there is a need to conclude from IInfo what is the new leaf (leaf *new* in the INSERT routine). For simplicity of presentation, we do not add this field, and abstractly refer to it in the code.

```

FIND(Key  $k$ ) : Leaf* {
1   Internal * $gp$ , * $p$ 
2   Leaf * $l$ 
3   Update  $pupdate, gpupdate$ 
4   while (TRUE) {
5        $\langle gp, p, l, pupdate, gpupdate \rangle := \text{SEARCH}(k)$ 
6       if  $gpupdate.state \neq \text{CLEAN}$  then HELP( $gpupdate$ )
7       else if  $pupdate.state \neq \text{CLEAN}$  then HELP( $pupdate$ )
8       else if  $gp \leftarrow update = gpupdate$  and  $p \leftarrow update = pupdate$  then {
9           if  $l \rightarrow key = k$  then return  $l$ 
10          else return  $\perp$ 
11      }
12  }
13 }

```

Figure 1: Solution 1: R-linearizability FIND routine

The correctness of the two solutions relies on the following argument. Once a process flags a node during operation Op with input key k (either INSERT or DELETE), then if this attempt to complete the operation eventually succeed (i.e., the marking is also successful in the case of DELETE), then any FIND(k) operation invokes from this point consider Op as if it is completed.

The suggested modification, although being simple and local, only guarantee the implementation satisfy R-linearizability. However, the problem of response being lost in case of a crush is not addressed. In general, the critical points in the code for recovery are the CAS primitives, as a crush right after applying CAS operation results the lost of the response, and in order to complete the operation the process needs to know the result of the CAS. In addition, because of the helping mechanism, a suspended DELETE operation which flagged a node, and yet to mark one, may be completed by other process in the future, and may not. Upon recovery, the process needs to distinguish between the two cases, in order to obtain the right response.

To address this issue, the expend the helping mechanism, so that the helping process needs also to update the info structure in case of a success. This is done by adding a boolean field to the Flag structure. This way, if a process crush along an operation Op , upon recovery it can check whether the operation was completed by some different process.

Before a process attempt to perform an operation, as it creates the Flagstructure op describing the operation and its affect on the data structure, the process stores op in a designated location (according to its id). Upon recovery, the process reads this location, and if the operation is not complete, then it retries to perform it, starting from the point of the first flagging (the first CAS). Otherwise, the operation was completed, and the response value is already known. Notice that there is a scenario in which process q recovers and observes an operation Op as not being complete, but just before it retries it, some other process complete the operation. We need to prove that even in such case, the operation will affect the data structure exactly once.

The supplied implementation does not specify what happens if a process crush outside of a BST routine.

```

FIND(Key k) : Leaf* {
14   Internal *gp, *p
15   Leaf *l
16   Update pupdate, gpupdate

17    $\langle gp, p, l, pupdate, gpupdate \rangle := \text{SEARCH}(k)$ 
18   if  $l \rightarrow key \neq k$  then {
19       if (pupdate.state = IFlag and pupdate.info attempt to add key k) then
20           return leaf with key k from pupdate.info
21       else return  $\perp$ 
22   }
23   if (pupdate.state = MARK and pupdate.info  $\leftarrow l \leftarrow key = k$ ) then return  $\perp$ 
24   if (gpupdate.state = DFlag and gpupdate.info  $\leftarrow l \leftarrow key = k$ ) then {
25       op := gpupdate.info
26       result := CAS(op  $\rightarrow p \rightarrow update$ , op  $\rightarrow pupdate$ ,  $\langle \text{MARK}, op \rangle$ )     $\triangleright$  mark CAS
27       if (result = op  $\rightarrow pupdate$  or result =  $\langle \text{MARK}, op \rangle$ ) then return  $\perp$      $\triangleright$  op  $\rightarrow p$  is successfully marked
28   }
29   return l
30 }

```

Figure 2: Solution 2: R-linearizability FIND routine

```

31 type Update {                 $\triangleright$  stored in one CAS word
32     {CLEAN, DFlag, IFlag, MARK} state
33     Flag *info
34 }
35 type Internal {               $\triangleright$  subtype of Node
36      $Key \cup \{\infty_1, \infty_2\}$  key
37     Update update
38     Node *left, *right
39 }
40 type Leaf {                   $\triangleright$  subtype of Node
41      $Key \cup \{\infty_1, \infty_2\}$  key
42 }
43 type IInfo {                  $\triangleright$  subtype of Flag
44     Internal *p, *newInternal
45     Leaf *l
46     boolean complete
47 }
48 type DInfo {                  $\triangleright$  subtype of Flag
49     Internal *gp, *p
50     Leaf *l
51     Update pupdate
52     boolean complete
53 }
 $\triangleright$  Initialization:
54 shared Internal *Root := pointer to new Internal node
    with key field  $\infty_2$ , update field  $\langle \text{CLEAN}, \perp \rangle$ , and
    pointers to new Leaf nodes with keys  $\infty_1$  and
     $\infty_2$ , respectively, as left and right fields.

```

Figure 3: Type definitions and initialization.

```

RECOVER() {
55     Flag *op = Announce[id]

56     if op of type IInfo then
57         if op ← complete = TRUE then return TRUE
58         else go to line 1
59     if op of type DInfo then
60         if op ← complete = TRUE then return TRUE
61         else go to line 1
62 }

```

Figure 4: RECOVER routine

```

63 SEARCH(Key  $k$ ) : (Internal*, Internal*, Leaf*, Update, Update) {
    ▷ Used by INSERT, DELETE and FIND to traverse a branch of the BST; satisfies following postconditions:
    ▷ (1)  $l$  points to a Leaf node and  $p$  points to an Internal node
    ▷ (2) Either  $p \rightarrow \text{left}$  has contained  $l$  (if  $k < p \rightarrow \text{key}$ ) or  $p \rightarrow \text{right}$  has contained  $l$  (if  $k \geq p \rightarrow \text{key}$ )
    ▷ (3)  $p \rightarrow \text{update}$  has contained  $pupdate$ 
    ▷ (4) if  $l \rightarrow \text{key} \neq \infty_1$ , then the following three statements hold:
    ▷ (4a)  $gp$  points to an Internal node
    ▷ (4b) either  $gp \rightarrow \text{left}$  has contained  $p$  (if  $k < gp \rightarrow \text{key}$ ) or  $gp \rightarrow \text{right}$  has contained  $p$  (if  $k \geq gp \rightarrow \text{key}$ )
    ▷ (4c)  $gp \rightarrow \text{update}$  has contained  $gpupdate$ 
64 Internal * $gp$ , * $p$ 
65 Node * $l := \text{Root}$ 
66 Update  $gpupdate, pupdate$  ▷ Each stores a copy of an update field
67 while  $l$  points to an internal node {
68      $gp := p$  ▷ Remember parent of  $p$ 
69      $p := l$  ▷ Remember parent of  $l$ 
70      $gpupdate := pupdate$  ▷ Remember update field of  $gp$ 
71      $pupdate := p \rightarrow \text{update}$  ▷ Remember update field of  $p$ 
72     if  $k < l \rightarrow \text{key}$  then  $l := p \rightarrow \text{left}$  else  $l := p \rightarrow \text{right}$  ▷ Move down to appropriate child
73 }
74 return  $\langle gp, p, l, pupdate, gpupdate \rangle$ 
75 }

76 FIND(Key  $k$ ) : Leaf* {
77     Leaf * $l$ 
78      $\langle -, -, l, -, - \rangle := \text{SEARCH}(k)$ 
79     if  $l \rightarrow \text{key} = k$  then return  $l$ 
80     else return  $\perp$ 
81 }

82 INSERT(Key  $k$ ) : boolean {
83     Internal * $p, *newInternal$ 
84     Leaf * $l, *newSibling$ 
85     Leaf * $new$  := pointer to a new Leaf node whose key field is  $k$ 
86     Update  $pupdate, result$ 
87     IInfo * $op$ 
88     while TRUE {
89          $\langle -, p, l, pupdate, - \rangle := \text{SEARCH}(k)$ 
90         if  $l \rightarrow \text{key} = k$  then return FALSE ▷ Cannot insert duplicate key
91         if  $pupdate.state \neq \text{CLEAN}$  then HELP( $pupdate$ ) ▷ Help the other operation
92         else {
93              $newSibling$  := pointer to a new Leaf whose key is  $l \rightarrow \text{key}$ 
94              $newInternal$  := pointer to a new Internal node with key field  $\max(k, l \rightarrow \text{key})$ ,
95                 update field  $\langle \text{CLEAN}, \perp \rangle$ , and with two child fields equal to  $new$  and  $newSibling$ 
96                 (the one with the smaller key is the left child)
97              $op$  := pointer to a new IInfo record containing  $\langle p, l, newInternal, \text{FALSE} \rangle$ 
98              $\text{Announce}[id] := op$ 
99              $result := \text{CAS}(p \rightarrow \text{update}, pupdate, \langle \text{IFlag}, op \rangle)$  ▷ iflag CAS
100             if  $result = pupdate$  or  $result = \langle \text{IFlag}, op \rangle$  then { ▷ The iflag CAS was successful
101                 HELPINSERT( $op$ ) ▷ Finish the insertion
102                 return TRUE
103             }
104             else HELP( $result$ ) ▷ The iflag CAS failed; help the operation that caused failure
105         }
106     }
107     if  $op \rightarrow \text{complete} = \text{TRUE}$  then
108         return TRUE
109     }

110 HELPINSERT(IInfo * $op$ ) {
111     ▷ Precondition:  $op$  points to an IInfo record (i.e., it is not  $\perp$ )
112     CAS-CHILD( $op \rightarrow p, op \rightarrow l, op \rightarrow newInternal$ ) ▷ ichild CAS
113     CAS( $op \rightarrow p \rightarrow \text{update}, \langle \text{IFlag}, op \rangle, \langle \text{CLEAN}, op \rangle$ ) ▷ iunflag CAS
114      $op \rightarrow \text{complete} := \text{TRUE}$  ▷ mark the operation as completed
115 }

```

Figure 5: Pseudocode for SEARCH, FIND and INSERT.

```

113 DELETE(Key  $k$ ) : boolean {
114   Internal * $gp$ , * $p$ 
115   Leaf * $l$ 
116   Update  $pupdate$ ,  $gpupdate$ ,  $result$ 
117   DInfo * $op$ 

118   while TRUE {
119      $\langle gp, p, l, pupdate, gpupdate \rangle := \text{SEARCH}(k)$ 
120     if  $l \rightarrow key \neq k$  then return FALSE ▷ Key  $k$  is not in the tree
121     if  $gpupdate.state \neq \text{CLEAN}$  then HELP( $gpupdate$ )
122     else if  $pupdate.state \neq \text{CLEAN}$  then HELP( $pupdate$ )
123     else { ▷ Try to flag  $gp$ 
124        $op :=$  pointer to a new DInfo record containing  $\langle gp, p, l, pupdate, \text{FALSE} \rangle$ 
125        $\text{Announce}[id] := op$ 
126        $result := \text{CAS}(gp \rightarrow update, gpupdate, \langle \text{DFlag}, op \rangle)$  ▷ dflag CAS
127       if  $result = gpupdate$  or  $result = \langle \text{DFlag}, op \rangle$  then { ▷ CAS successful
128         if HELPDELETE( $op$ ) then return TRUE ▷ Either finish deletion or unflag
129       }
130       else HELP( $result$ ) ▷ The dflag CAS failed; help the operation that caused the failure
131     }
132     if  $op \rightarrow complete = \text{TRUE}$  then
133       return TRUE
134   }
135 }

136 HELPDELETE(DInfo * $op$ ) : boolean {
137   ▷ Precondition:  $op$  points to a DInfo record (i.e., it is not  $\perp$ )
138   Update  $result$  ▷ Stores result of mark CAS
139    $result := \text{CAS}(op \rightarrow p \rightarrow update, op \rightarrow pupdate, \langle \text{MARK}, op \rangle)$  ▷ mark CAS
140   if  $result = op \rightarrow pupdate$  or  $result = \langle \text{MARK}, op \rangle$  then { ▷  $op \rightarrow p$  is successfully marked
141     HELPMARKED( $op$ ) ▷ Complete the deletion
142     return TRUE ▷ Tell DELETE routine it is done
143   }
144   else { ▷ The mark CAS failed
145     HELP( $result$ ) ▷ Help operation that caused failure
146      $\text{CAS}(op \rightarrow gp \rightarrow update, \langle \text{DFlag}, op \rangle, \langle \text{CLEAN}, op \rangle)$  ▷ backtrack CAS
147     return FALSE ▷ Tell DELETE routine to try again
148   }

149 HELPMARKED(DInfo * $op$ ) {
150   ▷ Precondition:  $op$  points to a DInfo record (i.e., it is not  $\perp$ )
151   Node * $other$ 
152   ▷ Set  $other$  to point to the sibling of the node to which  $op \rightarrow l$  points
153   if  $op \rightarrow p \rightarrow right = op \rightarrow l$  then  $other := op \rightarrow p \rightarrow left$  else  $other := op \rightarrow p \rightarrow right$ 
154   ▷ Splice the node to which  $op \rightarrow p$  points out of the tree, replacing it by  $other$ 
155    $\text{CAS-CHILD}(op \rightarrow gp, op \rightarrow p, other)$  ▷ dchild CAS
156    $\text{CAS}(op \rightarrow gp \rightarrow update, \langle \text{DFlag}, op \rangle, \langle \text{CLEAN}, op \rangle)$  ▷ dunflag CAS
157    $op \rightarrow complete := \text{TRUE}$  ▷ mark the operation as completed
158 }

159 HELP(Update  $u$ ) { ▷ General-purpose helping routine
160   ▷ Precondition:  $u$  has been stored in the  $update$  field of some internal node
161   if  $u.state = \text{IFlag}$  then HELPINSERT( $u.info$ )
162   else if  $u.state = \text{MARK}$  then HELPMARKED( $u.info$ )
163   else if  $u.state = \text{DFlag}$  then HELPDELETE( $u.info$ )
164 }

165 CAS-CHILD(Internal * $parent$ , Node * $old$ , Node * $new$ ) {
166   ▷ Precondition:  $parent$  points to an Internal node and  $new$  points to a Node (i.e., neither is  $\perp$ )
167   ▷ This routine tries to change one of the child fields of the node that  $parent$  points to from  $old$  to  $new$ .
168   if  $new \rightarrow key < parent \rightarrow key$  then
169      $\text{CAS}(parent \rightarrow left, old, new)$ 
170   else
171      $\text{CAS}(parent \rightarrow right, old, new)$ 
172 }

```

Figure 6: Pseudocode for DELETE and some auxiliary routines.